

<http://127.0.0.1:8000/>

serv-ing-stack

wk2

OAS 3.1

/openapi.json

default

GET

/health

Health

Liveness and readiness.

Contract: returns 200 with {"status": "ok", "model": ""} once the model is loaded. Kubernetes probes (week 4) and the agentic client's retry logic (weeks 4 to 6) call this. It must be cheap and must not run the model.

Parameters

No parameters

Cancel

Execute

Clear

Responses

Curl

```
curl -X 'GET' \
  'http://127.0.0.1:8000/health' \
  -H 'accept: application/json'
```

Request URL

```
http://127.0.0.1:8000/health
```

Server response

Code	Details
200	Response body <pre>{ "status": "ok", "model": "Queen/Queen2.5-0.5B-Instruct" }</pre> Download Response headers <pre>content-length: 52 content-type: application/json date: Mon, 24 Aug 2026 15:37:46 GMT server: uvicorn</pre>

Responses

Code	Description	Links
200	Successful Response Media type application/json Controls Accept header. Example Value Schema <pre>{ "status": "ok", "model": "string" }</pre>	

 No links |

```
(aidc-env) shurooq@ShurooqAljohani:/mnt/c/Users/shoro/Documents/aidc-bootcamp/serving-stack$ python3 app/client_test.py
reply: Three primary colors are red, yellow, and blue.
finish_reason: stop
usage: CompletionUsage(completion_tokens=12, prompt_tokens=24, total_tokens=36, completion_tokens_details=None, prompt_t
okens_details=None)
(aidc-env) shurooq@ShurooqAljohani:/mnt/c/Users/shoro/Documents/aidc-bootcamp/serving-stack$
```

```

NameError: name 'tokenizer' is not defined
WARNING: WatchFiles detected changes in 'main.py'. Reloading...
INFO: Shutting down
INFO: Waiting for application shutdown.
INFO: Application shutdown complete.
INFO: Finished server process [22636]
INFO: Started server process [22787]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: 127.0.0.1:38216 - "POST /v1/chat/completions HTTP/1.1" 200 OK
INFO: 127.0.0.1:44712 - "GET /docs HTTP/1.1" 200 OK
INFO: 127.0.0.1:44712 - "GET /openapi.json HTTP/1.1" 200 OK
INFO: 127.0.0.1:44718 - "POST /v1/chat/completions HTTP/1.1" 200 OK
INFO: 127.0.0.1:46174 - "POST /v1/chat/completions HTTP/1.1" 200 OK

```

The screenshot shows a REST client interface with the following details:

- Method:** POST
- URL:** /v1/chat/completions
- Request body:** A JSON object with a model, messages, and a stream flag.
- Response:** A 200 OK status with a JSON body containing a chat completion response, including a message and token counts.
- Response headers:** Content-Length, Content-Type, Date, and Server.
- Validation:** A 422 Validation Error is shown for a different request, indicating a schema mismatch.

```

./verify-env-report.json
(aidc-env) shurooq@ShurooqAljohani:/mnt/c/Users/shoro/Documents/aidc-bootcamp/serving-stack$ nano verify.py
(aidc-env) shurooq@ShurooqAljohani:/mnt/c/Users/shoro/Documents/aidc-bootcamp/serving-stack$ python3 verify.py
GREEN CHECK: PASS
(aidc-env) shurooq@ShurooqAljohani:/mnt/c/Users/shoro/Documents/aidc-bootcamp/serving-stack$

```